

Regression Task: Predicting Second Hand Car Prices

Buğracan Tanrıverdi

1. Task

The task we have given is to predict the price of second hand cars. To do this, we will use this data set [Kaggle](#). Since price is a continuous variable we will use regression for this task.

2. Description of the Data Set

First of all, we must import the data set.

```
cars <- read.csv("cars.csv")
cars <- cars[,2:12]
head(cars)
```

	on.road.old	on.road.now	years	km	rating	condition	economy	top.speed	hp
1	535651	798186	3	78945	1	2	14	177	73
2	591911	861056	6	117220	5	9	9	148	74
3	686990	770762	2	132538	2	8	15	181	53
4	573999	722381	4	101065	4	3	11	197	54
5	691388	811335	6	61559	3	9	12	160	53
6	650007	844846	6	148846	2	9	13	138	61

	torque	current.price
1	123	351318.0
2	95	285001.5
3	97	215386.0
4	116	244295.5
5	105	531114.5
6	109	177933.5

The data set consists of 12 columns (one of them is just ids, so we deleted that column) and 1000 rows. Since `current.price` is our *target* this means we have 10 *features*. Also all of our *features* seems like they are all `numeric`.

3. Model Training

3.1 Splitting the Data Set

Before training a regression model we should split the data set into two parts; `train` and `test`.

```
set.seed(1)
index <- sample(1 : nrow(cars), round(nrow(cars) * 0.80))
train <- cars[index, ]
test  <- cars[-index, ]
```

3.2 Training

We will train the model on the `train` subset of our data set using the `lm` function.

```
lm.model <- lm(current.price ~ ., data = train)
summary(lm.model)
```

Call:

```
lm(formula = current.price ~ ., data = train)
```

Residuals:

Min	1Q	Median	3Q	Max
-12415	-7382	-1864	5132	21921

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-1.507e+04	6.722e+03	-2.242	0.0252	*
on.road.old	5.057e-01	5.231e-03	96.685	<2e-16	***
on.road.now	5.000e-01	5.411e-03	92.401	<2e-16	***
years	-1.584e+03	1.787e+02	-8.865	<2e-16	***
km	-3.991e+00	1.058e-02	-377.268	<2e-16	***
rating	1.385e+02	2.187e+02	0.633	0.5268	
condition	4.531e+03	1.088e+02	41.646	<2e-16	***
economy	6.017e+01	1.385e+02	0.434	0.6642	

```
top.speed  -1.323e+01  1.603e+01  -0.825   0.4096
hp          1.478e+01  1.508e+01   0.980   0.3275
torque      1.611e+01  1.476e+01   1.091   0.2755
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 8660 on 789 degrees of freedom
```

```
Multiple R-squared:  0.9953,    Adjusted R-squared:  0.9953
```

```
F-statistic: 1.683e+04 on 10 and 789 DF,  p-value: < 2.2e-16
```

It seems that the only significant features are `on.road.old`, `on.road.new`, `years`, `km`, `condition`.

3.3 Prediction

We can easily calculate the predicted prices using the `predict` function.

```
predicted.price <- predict(lm.model, test)
head(predicted.price)
```

```
      18      23      26      32      38      46
281916.2 333030.9 379000.9 464309.6 429237.8 370423.1
```

4. Performance

In order to measure the performance of the regression model that we have trained we will use the *root mean squared error (RMSE)* metric, since it is easy to calculate and gives the error in the same unit as our *target*.

```
error <- test$current.price - predicted.price
rmse.model <- sqrt(mean(error ^ 2))
rmse.model
```

```
[1] 9091.345
```

Intuitively this means that on average our predictions differ 9000 from the actual price of the cars. Since the average price of a car in our data set is 308520.2 this error looks acceptable.

5. Overfitting

We must also check if there is an overfitting problem.

```
rmse.test <- sqrt(mean(lm.model$residuals^2))
rmse.model - rmse.test
```

```
[1] 491.4751
```

Since the difference between the model and test RMSE is a large positive number, most of the learning comes from the model. We can conclude that we don't have any overfitting problems.

6. Predicting New Observations

We can also create our own observations to predict. We will create 5 new observations.

```
new.observations
```

	on.road.now	on.road.old	years	km	rating	condition	economy	top.speed	hp
1	700000	550000	4	90000	1	8	7	150	70
2	750000	400000	6	60000	3	7	10	170	80
3	750000	500000	8	100000	5	6	11	140	100
4	800000	650000	7	70000	5	8	12	150	60
5	900000	600000	3	80000	4	5	9	180	50

torque

1	90
2	80
3	80
4	110
5	100

Now we will predict the price of each of these cars using our model.

```
predicted.price.new <- predict(lm.model, new.observations)
predicted.price.new
```

```
      1      2      3      4      5
284864.2 346218.1 230482.2 461534.8 438060.3
```